

A `pre_main` function that runs before static initialization that precedes `main`

Document #: D9999
Date: 2014-12-10
Project: Programming Language C++
Evolution Working Group
Reply-to: Daniel Gutson
<daniel.gutson@tallertechnologies.com>
Christopher Kormanyos
<e_float@yahoo.com>

1 Motivation

Overview

Many systems require user-specific initialization. In particular, deeply embedded environments may have user-specific initialization that should be executed as soon as possible following power-on-reset of the microcontroller. Examples thereof include initialization of I/O ports, watchdog timers, clock systems, *etc.* For these environments, postponing user-specific initialization until `main` may detract from quality of design.

2 Proposed Solution

This document proposes changes in the standard text reflecting the specification of a new `pre_main` function that is called prior to static initialization that precedes `main`.

Modification to the standard text

1. [basic.start.main] Add a new paragraph §3.6.1p6

The `pre_main` function shall be called prior to *static initialization* (§3.6.2) that precedes `main`. The linkage (3.5) of `pre_main` is implementation-defined. The `pre_main` function is parameter-free. The return value of `pre_main` is `void`. The `pre_main` function lacks side-effects. The mechanism ensuring that `pre_main` lacks side-effects is implementation-defined.

3 Existing mechanisms and workarounds

We will now investigate existing mechanisms and workarounds that implement a `pre_main`-style mechanism.

Custom startup code

Many embedded systems developers write custom startup code that supports a kind of `pre_main` function. For example, we will now look at sample code showing how an implementation of startup code might potentially provide a call mechanism for `pre_main`.

```
void __my_startup()
{
    // Load the sreg register.
    asm volatile("eor r1, r1");
    asm volatile("out 0x3F, r1");

    // Set the stack pointer.
    asm volatile("ldi r28, lo8(__initial_stack_pointer)");
    asm volatile("ldi r29, hi8(__initial_stack_pointer)");

    // Load the stack pointer high and low registers.
    asm volatile("out 0x3E, r29");
    asm volatile("out 0x3D, r28");

    // A potential call mechanism for pre_main.
    pre_main();

    // Initialize statics from ROM to RAM.
    // Initialize default-initialized static RAM.
    crt::init_ram();

    // Call all ctor initializations.
    crt::init_ctors();

    // Call main (and never return).
    asm volatile("call main");

    // Catch an unexpected return from main.
    for(;;)
    {
        // Replace with a loud error if desired.
        mcal::wdg::secure::trigger();
    }
}
```

This example has been taken from a low-level initialization sequence for a popular 8-bit microcontroller. The code has been compiled and tested with GCC 4.8.1. After setting a CPU register, the stack pointer is initialized. Immediately following stack setup, `pre_main` is called. Note that `pre_main` is called prior to static initialization.

Commercially available microcontroller compilers

Some commercially available microcontroller compilers provide a custom hook (in the sense of `pre_main`) that is called before static initialization that precedes `main`. The IAR Systems C/C++ compiler, for instance, uses an implementation-specific function called `__low_level_init` for this purpose. The user is responsible for the content (if any) of `__low_level_init`.

4 Future Work

The motivation and justification for a potential `pre_main` is analogous in C and C++. Specifying `pre_main` could potentially be addressed in WG14 as well as WG21.

Along these lines, do we need two versions of `pre_main`? In particular,

```
void      ::pre_main(void); // Intended for C
void std::pre_main();      // Intended for C++
```

What is the proper name of a potential `pre_main`? Is `pre_init` a better name because it more clearly reflects when the function is called?

Despite the proposal that `pre_main` is to lack side-effects, it could be beneficial to allow `pre_main` to initialize certain *clearly identifiable* non-local variables having static storage duration. Embedded systems tool chains for C/C++ typically provide special linker sections with implementation-specific names such as `.noinit`, `.noclear`, *etc.* These are meant to store non-local variables having static storage duration that are not intended to undergo static initialization. Attributes such as `[[noclear]]` or `[[noinit]]` could be used to clearly identify these.

Consider, for example, the `reset_reason` in the following code.

```
typedef enum enum_reset_reason
{
    power_on_reset ,
    watchdog_reset ,
    software_reset
}
reset_reason_type;

[[noclear]] reset_reason_type reset_reason;
```

Here, `reset_reason` is intended to be initialized by `pre_main` in the application, not via conventional static initialization.

5 Discussion

TBD: Summarize the discussion regarding when `pre_main` should be called and why.

TBD: Summarize the discussion regarding the dangers of offering an open user-interface that precedes `main`. Will users run into inordinate amounts of trouble with this proposed interface?

6 Acknowledgments

TBD: Acknowledge the participants.

7 References

TBD: Any references?